


Automated Musical Tempo Estimation with Neural-Network-Based Onset Detection

Nathan Stephenson

What? Why?

- Beats per minute (BPM) of music
 - The tempo of a song
 - Synchronizing anything to music
 - Statistics, analyzing tempo and correlating it to genre, artist, etc.
- Can't you do it yourself? Google it?
 - Yes, but imagine doing trial and error for 1,000 songs
 - The closer we get to 100% accuracy, the more times we don't have to manually adjust BPM
- Focus on music with same BPM throughout



<https://www.flickr.com/photos/48423254@N00/12294167> CC BY 2.0 Paco from Badajoz, España

Move Your Feet by Junior Senior is in the key of A Minor, 118 BPM. This track was released in 2002.

www.notediscover.com > song > junior-senior-move-your...

Key & BPM/Tempo of Move Your Feet by Junior Senior

getsongbpm.com > ... > D-D-Don't Don't Stop the Beat

BPM for Move Your Feet (Junior Senior) - GetSongBPM

Move Your Feet is played at 117 Beats Per Minute (Moderato),

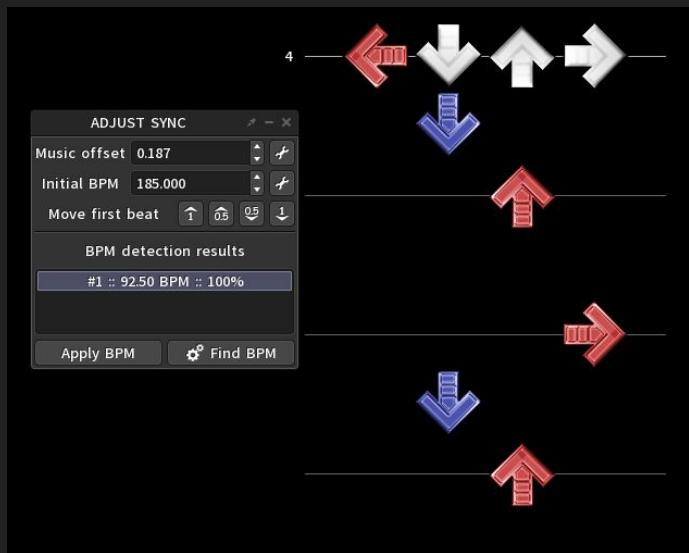
songbpm.com > move-your-feet

BPM for "Move Your Feet" by Junior Senior | SongBPM

"Move Your Feet" by Junior Senior has a tempo of 119 BPM.

The technology

- Bram van de Wetering
 - ArrowVortex
 - Paper called *Non-Causal Beat Tracking for Rhythm Games*



Method	Song	BPMs in order of highest to lowest confidence	Actual BPM
van de Wetering (ArrowVortex)	Bruno Mars - Perm (pop music with integer BPM)	124	124
Böck et. al (madmom)	Bruno Mars - Perm	61.86 (123.72), 125, 0.52	124
van de Wetering (ArrowVortex)	Junior Senior - Move Your Feet (pop music with unusual BPM)	118.868, 178.33	118.868*
Böck et. al (madmom)	Junior Senior - Move Your Feet	59.41 (118.82), 117.65	118.868*
van de Wetering (ArrowVortex)	A-One - STAR LINER (upbeat eurobeat with lots of percussion)	162, 108, 129.60	162
Böck et. al (madmom)	A-One - STAR LINER	162.16, 40.54, 0.62	162
van de Wetering (ArrowVortex)	Doobie Brothers - Toulouse Street (soft folk rock)	179.80, 179.32, 180.16	90 (180)**
Böck et. al (madmom)	Doobie Brothers - Toulouse Street (soft folk rock)	89.55 (179.1), 44.78, 0.53	90 (180)**
Average percent error (from closest BPM detected) van de Wetering (ArrowVortex): 0.018% Böck et. al (madmom): 0.743%			

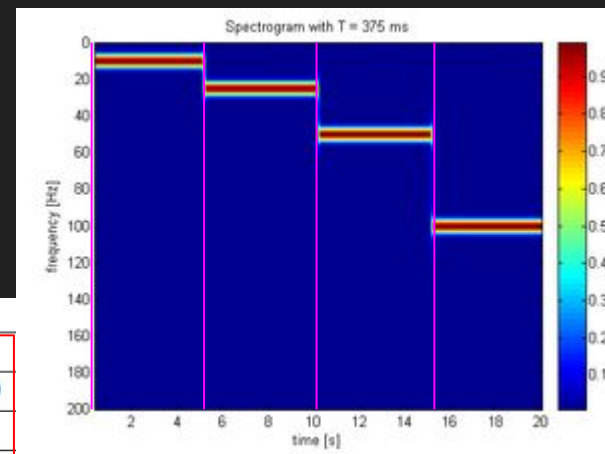
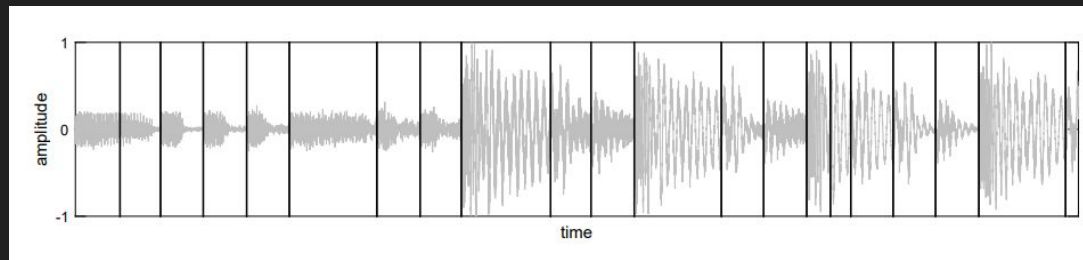
Steps

Given an audio file (30 seconds or more, preferably)

1. Detect onsets
2. Fill coarse intervals
3. Refine intervals and choose the best
4. Calculate BPM

Onset detection

- Detect onsets
 - The start of a sound or musical note
 - Uses frequency and phase information
 - Allows us to view audio as frequency over time
- Spectral Flux determined to be the best method



Method	HFC	CD	PB	SD	KL	MKL	SF
Accurate	100/100	100/100	83/100	98/100	100/100	97/100	100/100
SSE	0.008	0.007	0.214	0.012	0.005	0.045	0.003


Table 1: Results of the onset detection functions evaluation, showing accuracy and sum of squared errors (SSE).

https://commons.wikimedia.org/wiki/File:STFT_colored_spectrogram_375ms.png CC BY-SA 3.0

Alessio Damato

Coarse intervals

- Choose a range of tempo values
 - 89-205 BPM
- Interval in samples = sample rate (44100 Hz) * 60 seconds / n beats per minute
 - Number of samples per beat
 - 60 BPM -> 1 beat per second -> 44100 Hz
- Loop over the interval in increments of 10, making it a “coarse” runthrough



from *Non-causal Beat Tracking for Rhythm Games*


Refine intervals

- Go back and find our best matches and loop over them with a delta of 1

Calculate BPM

- Interval to BPM
 - $\text{sample rate (Hz)} * 60 \text{ seconds} / (\text{interval} * 1 \text{ minute})$

...and then we get the BPM! (hopefully!)



from *Non-causal Beat Tracking for Rhythm Games*

My work


- Author sent part of the ArrowVortex code
- Heavily modified it to be standalone
 - Fixed it up, but it gave me completely different values
- Reverse engineering with Ghidra and x64dbg
 - “Decompiles” the machine language into (barely) readable code!
 - x64dbg, real-time debugger which lets me find onsets in memory (RAM)
- Complex Domain used, not Spectral Flux
- The onsets my code generates were still notably different

```
aubio_onset_t* o = new_aubio_onset("complex", window, hop_size, samplerate);
do {
    for (int i = 0; i < hop_size; i++) {
        // if (total + i >= numFrames) PrintOut("exceeding numFrames")
        in->data[i] = (total+i < numFrames) ? samples[total+i] : 0;
    }

    aubio_onset_do(o, in, out);

    // if (out->data[0] > 0) // if there is an onset
    if (out->data[0] > 0 && aubio_onset_get_last(o) >= 0) {
        printf("%d\n", aubio_onset_get_last(o));
        onsets->append(Onset(aubio_onset_get_last(o)));
    }

    total += hop_size;
} while (total < numFrames);
```



```
_in = *(uint **) (unaff_EBP + -0x18);
aubio_onset_do((int *)onset_finder_in, (int*)_out,
_out = *(int **) (unaff_EBP + 0x14);
if ((0.00000000 < *(float *)_out[1]) &&
    (-1 < (int) (uint *) ((int)onset_finder[0xb] - (int)onset_finder[7]))) {
    *(uint **) (unaff_EBP + -0x28) = (uint *) ((int)onset_finder[0xb] - (int)onset_finder[7]);
    *(undefined8 *) (unaff_EBP + -0x20) = 0x3ff0000000000000;
    probably_append(*(void **) (unaff_EBP + 0x18), (undefined4 *) (unaff_EBP + -0x28));
    _out = *(int **) (unaff_EBP + 0x14);
}
pUVar3 = (undefined4 *) ((int *) (unaff_EBP + -0x10) + 0x400);
pIVar1 = (int *) (unaff_EBP + -0x14);
*piVar1 = *pIVar1 + -1;
*(undefined4 **) (unaff_EBP + -0x10) = pUVar3;
```

(gdb) x/32xh &onsets->begin()

0x7fffffffedaa0:	0xba70	0x0846	0x0000	0x0000	0xb560	0x0844	0x0000	0x0000
0x7fffffffedab0:	0xaea0	0x0846	0x0000	0x0000	0xee1c	0xff1d	0x7fff	0xcc00
0x7fffffffedac0:	0xc160	0x0844	0xea58	0x4186	0xb3e0	0x0844	0x0000	0x0000
0x7fffffffedad0:	0xa5a0	0x0844	0x0000	0x0000	0x9a80	0x0846	0x0000	0x0000

Address	Index	ASCI
77481000	18 00 38 00	80 C1 48 77
77481010	28 00 2A 00	84 C1 48 77
77481020	1E 20 00 00	1A 00 02 00
77481030	18 00 1A 00	D4 C0 48 77
77481040	30 00 32 00	7C C0 48 77
77481050	20 00 22 00	2C C0 48 77
77481060	10 00 12 00	E8 BE 48 77
77481070	18 00 00 00	84 17 48 77
77481080	00 00 00 00	00 00 00 00
77481090	7C 17 48 77	40 00 00 00
774810A0	18 00 00 00	00 00 00 00
774810B0	14 00 16 00	28 7C 48 77
774810C0	0E 00 00 00	00 00 00 00
774810D0	0E 00 00 00	00 00 00 00
774810E0	08 00 0A 00	6C 78 48 77
774810F0	06 00 08 00	20 7D 48 77
77481100	06 00 08 00	E8 7D 48 77
77481110	18 00 00 00	00 00 00 00
77481120	00 00 00 00	57 14 01 E2
77481130	E8 03 F0 06	46 15 C5 43
77481140	9A 8B 13 35	96 50 BD 4F
77481150	06 00 01 00	40 74 48 77
77481160	89 33 41 44	FE 06 06 30
77481170	24 74 48 77	03 00 00 00
77481180	95 8B 83 7D	76 6C 67 1F
77481190	04 00 00 00	12 7A 0F E8
774811A0	50 A1 5A 9A	0A 00 00 00
774811B0	40 7C 48 77	02 00 04 00
774811C0	24 7D 48 77	00 C8 40 77

Fixing my work

- I received more code to reference
 - Two lines of incorrect code broke everything!
 - Accidentally used a single byte instead of an integer
 - Maximum interval was 255 but I needed it to be ~17,000
- Much better!

Song	van de Wetering (<u>ArrowVortex</u>)	My current code	Böck et. al (madmom)	Actual BPM
Bruno Mars - Perm	124	124	61.86 (123.72)	124
Junior Senior - Move Your Feet	118.868	118.873	59.41 (118.82)	118.88*
A-One - STAR LINER	162	162	162.16	162
Doobie Brothers - Toulouse Street	179.80	178.11, 181.61	89.55 (179.1)	90 (180)**
<p>* This tempo is accurate enough that from beginning to end the beats do not noticeably deviate from the audio (roughly $\pm 10\text{ms}$).</p> <p>** This tempo is variable and an average integer tempo is suggested.</p> <p>*** BPM values with parentheses are the originals multiplied by 2 to match other values, as multiplying or dividing BPM values by powers of 2 do not change the synchronization of the beats.</p>				

Comparing onset methods

- Improving the algorithm and surpassing the state-of-the-art!
- Spotify has 30-second previews of songs
- Use songs with a constant BPM
 - Top Spotify tracks from Billboard's Top Artists of the 2010s list + a few other artists
 - Fixed BPM uncommon until the mid-'70s or '80s
 - Mostly pop and rap music being tested
 - Most songs should have a fixed tempo
- Neural-network-based algorithms found in madmom Python library
 - BLSTM, CNN, LL (Real-Time Online RNN)
 - Sebastian Böck
- Other algorithms come from the aubio Python library
 - **Complex Domain**, Energy-Based Distance, High-Frequency Content, Modified Kullback-Leibler, Kullback-Leibler, Phase-Based, Spectral Difference, **Spectral Flux**
- Testing Spotify's own tempo detection algorithm as well

Song	<i>BLSTM</i>	Complex	<i>CNN</i>	Energy	HFC	<i>LL (RNN)</i>	MKL	KL	PB	SpecDiff	SpecFlux	Actual
Bruno Mars - Perm	124	124	124	124	124	124	124	124	124	124	124	124
Junior Senior - Move Your Feet	118.879	118.879	118.879	118.879	118.879	118.879	118.879	118.879	118.879	118.879	118.879	118.879*
A-One - STAR LINER	162	162	162	162	162	162	162	162	189.596	162	162	162
Doobie Brothers - Toulouse Street	179.402	179.378	179.817	179.841	179.378	179.817	179.366	179.817	179.390	179.817	179.329	90 (180)**

* This tempo is accurate enough that from beginning to end the beats do not noticeably deviate from the audio (roughly $\pm 10\text{ms}$).

** This tempo is variable and an average integer tempo is suggested.

Method (italicized = machine learning, bolded = used in van de Wetering/ArrowVortex)	Mean squared error (using mode as “ground truth”)	Significant error count (>= 0.05 BPM)	Insignificant error count	Half-BPM detections
<i>Bidirectional LSTM</i>	13.831	69/772 (8.94%)	198	2
Complex Domain	17.367	47/772 (6.09%)	123	1
<i>CNN</i>	8.435	56/772 (7.25%)	196	3
Energy-Based Distance	54.111	129/772 (16.7%)	147	10
High-Frequency Content	44.098	67/772 (8.68%)	120	2
<i>LL (Online RNN)</i>	14.171	40/772 (5.18%)	151	0
Modified Kullback-Leibler	16.197	54/772 (6.99%)	115	2
Kullback-Leibler	27.504	75/772 (9.72%)	141	2
Phase-Based	262.092	287/772 (37.18%)	158	33
Spectral Difference	36.695	61/772 (7.90%)	134	6
Spectral Flux	33.334	47/772 (6.09%)	123	1
Spotify	116.614	388/771 (50.32%)	375	264

So, which is the best?

- Not Spotify!
 - The BPM values you can easily Google search are likely based on Spotify's data
- Precision
 - **CNN**, BLSTM, LL
- Reliability
 - **LL**, **Complex Domain**, Spectral Flux
- Performance
 - Machine learning methods are slower
 - **Complex Domain** and **Spectral Flux** seem to stand out still as the best non-ML-based algorithms

What next?

- Tempo estimation is not perfect
- Testing a variety of genres
 - Collecting accuracy by genre
- Focus on music with less defined onsets
 - Classical music, softer jazz, ambient
 - Music without drums in general

Thanks for listening!